onboard ai

A FIELD REPORT FROM 50+ AI ASSESSMENTS · HEALTHCARE AI GOVERNANCE

Trust isn't a value.
It's an artifact.

What 1,452 control evaluations across eight health systems taught us about the difference between an AI review that reassures you and one that protects you.

50+ AI tools assessed	**8** health systems	**1,452** control evaluations	**1.6×** more risk surfaced when evidence is required

The finding, in one sentence.

Across the same tools and the same questionnaires, the questions that required evidence surfaced 1.6× the risk of the questions that accepted a description.

Ask a developer to describe how their tool was validated and the answer comes easily. Ask them to produce the artifact behind the claim — the validation report, the subgroup performance table, the monitoring plan — and the picture changes. Same developers, same tools; the only variable was what we asked them to prove.

For a CIO, that gap is not academic. It is the difference between a governance program that generates green dashboards and one that generates defensible decisions — and it shows up in the two places that matter most: the committee hours you spend now, and the liability you carry later.

1 · DESCRIPTION IS NOT EVIDENCE

A questionnaire can ask a vendor to “describe how the model was validated,” or it can ask them to “attach the validation results.” Those feel similar. They are not. When we hold the tool and the reviewer constant and change only the ask, the failure rate roughly doubles — from about one control in five to two in five. The evidence request is simply a stronger instrument; it catches what a description lets pass.

This is the core insight of the report: an AI assessment measures the questions you ask as much as it measures the tool in front of you. A review that never asks for evidence will tend to come back clean — not because the AI is sound, but because nothing in the process was built to find out.

> *“The review that asks for less isn't cheaper. You just pay for it later.”*

2 · TIMING DECIDES WHOSE JOB THE FIX BECOMES

Every risk a review surfaces produces a mitigation, and every mitigation lands with an owner. A median AI tool carries roughly 30 — about 12 that belong to the developer as contract commitments, and 18 that belong to the implementer's own clinical, IT, and security teams. The implementer's share is theirs no matter what. The developer's share is only assignable while the contract is open.

Surface a risk before signature and the fix can be written into the agreement. Surface it after go-live and there is no counterparty on the hook — the work doesn't disappear, it defaults to you. Same evidence either way; only the timing, and the owner, change. The good developers would rather be asked early, too.

3 · THE EVIDENCE ALMOST NOBODY CAN PRODUCE

The Joint Commission and CHAI guidance on the responsible use of AI (September 2025) names risk and bias assessment as one of seven foundational elements — asking health systems to document "performance across diverse populations," at procurement and again in use. It is a reasonable ask, and the single hardest thing to obtain.

When we asked for exactly that — a performance breakdown by race, ethnicity, sex, and age — four of thirty-six tools could produce it: the worst-covered requirement in everything we have assessed. You only discover the gap if your questionnaire asks for the artifact rather than the assurance.

> JOINT COMMISSION + CHAI GUIDANCE · RISK & BIAS ASSESSMENT
>
> **"Performance across diverse populations — assessed at procurement, and again in use."**
>
> Guidance is voluntary today. A voluntary Joint Commission certification built on the CHAI playbooks opened in June 2026.

4 · FROM PRINCIPLES TO CERTIFICATION, IN NINE MONTHS

SEP 2025	MAY 2026	JUN 2026
Guidance	**CHAI playbooks**	**JC certification**
7 elements · non-binding	baseline controls	voluntary · 22,000+ eligible

The pace matters. Meanwhile the AI keeps arriving — tools entering review roughly doubled quarter over quarter. The portfolio you clear this quarter is the one a certification would examine next.

None of this requires alarm. It requires an inventory: every AI tool, where it is deployed, the risks it carries, and the evidence behind each one — in the shape a surveyor, a board, or a clinician would accept. Collected at intake, that record is nearly free. Reconstructed after the fact, it is a project.

onboard ai

TAKE THIS TO YOUR NEXT AI GOVERNANCE COMMITTEE

Three questions worth asking about your own AI review

1. How many of your assessment controls require an artifact — not a description? If the answer is zero, a clean result tells you very little.
2. Which of your surfaced risks are the developer's to fix, and are they written down while the contract is still open?
3. Could you produce a subgroup performance breakdown for your deployed AI tools today — in the form the guidance asks for?

WHAT WE'D BUILD WITH YOU

One inventory of your AI — vendor products, EMR features, and internal builds alike — showing where each tool is deployed, the risks it carries, the owner of every mitigation, and the evidence behind each one. Kept current, and kept in the shape a board or a surveyor would accept. It is the difference between believing your AI is safe and being able to show it.

We'd rather co-develop with you than sell to you.

These figures come from de-identified assessments across eight health systems. If you're working through the same questions, we're always glad to compare notes.

Compare notes → on-board.ai