

OLIVER WYMAN
A MARSH BUSINESS

GLOBAL FINANCIAL MARKET INFRASTRUCTURE REPORT 2026

Back to the future

Key trends and strategic imperatives
for tomorrow's industry leaders



Contents

Executive summary	3
Industry outlook	6
Deep dive: Artificial intelligence	14
Deep dive: New markets	27
Conclusion	42

Executive summary

Between 2020 and 2025, the Financial Market Infrastructure (FMI) sector grew at an impressive compound annual growth rate (CAGR) of about 15%, expanding to \$131 billion revenue by the end of the five years. The surge was driven by volatility and secular growth across most asset classes, the deepening of capital markets, increased retail participation, financialization of physical asset classes, and continued demand for risk management. The cash generated from these highly profitable, deep-moat core infrastructure businesses funded inorganic expansion into higher-multiple and recurring businesses in data and technology. The sector cumulatively deployed more than \$100 billion for acquisitions, \$88 billion of which was spent on data and technology business line additions, which now account for 30% of sector revenue versus 21% five years ago. This has created more recurring revenue, broadened customer reach, and increased cross-sell into buy-side, sell-side, and corporate clients.

The rise of enterprise AI is now stress-testing which of those acquisitions are truly synergistic. Diversification into strategically and operationally synergistic adjacencies should prove valuable and defensible. Less connected extensions, disproportionately concentrated in commoditized data, analytics and workflow acquired in the last cycle, will face greater scrutiny and risk becoming challenged or a drag on FMI group value.

Enterprise AI is now creating both revenue opportunity and threat across FMI data and technology business lines. Exposure is highest in replicable data, analytics and workflow components that lack proprietary content, deep embedding, network effects, regulatory relevance or system-of-work/source-of-origin status. We estimate around 9% of industry revenue (close to \$12 billion) is at direct risk, with an additional 16% (or \$21 billion) at risk without active mitigation. The ability to retain revenue is strongest where data and workflows are proprietary, authoritative, mission-critical, embedded in client processes, or tied to core infrastructure. The priority for FMI is to link vulnerable components into end-to-end solutions that combine commoditized functionality with moated data, workflow, network, and control points.

Yet AI could also unlock around \$9 billion of incremental revenue, on top of underlying sector growth. This would be through three levers: new AI-enabled products, pricing uplift tied to measurable outcomes, and higher product usage from agentic AI use.

Enterprise AI also creates a material cost and operating leverage opportunity for the sector.

We estimate that AI can deliver more than a 20% cost reduction, raising the sector's margins on earnings before interest, taxes, depreciation, and amortization (EBITDA) by around nine percentage points from 54% to 63% for the sector.

The largest opportunities sit in software engineering and quality assurance, operations, risk, compliance, surveillance, and corporate functions, where AI can compress cycle times, improve first-time-right rates, reduce cost-to-serve and increase reuse across products and geographies. The opportunity is harder to capture in regulated infrastructure businesses than in less regulated data and technology business lines, given higher requirements for resilience, auditability, explainability and model risk control.

Realizing these benefits at scale requires an AI factory approach, industrializing use cases via reusable components, shared data layers, and a value control tower. Scaling AI safely will require FMIs to make governance part of the operating model, with a top-of-house framework for risk appetite, use-case classification, approval paths, model risk management, and clear accountability across the lines of defense.

Meanwhile, the impact of tokenization is disrupting core infrastructure rails while also expanding the perimeter of market infrastructure. Tokenization creates both an offensive and defensive imperative for FMIs, with as much as 35% of today's revenue (around \$45 billion) potentially migrating to new rails and around 5% (\$7 billion) structurally eliminated altogether. The downside is most acute where post-trade and trade franchises potentially lose their market-of-record status.

Yet broader asset-universe expansion could lift the sector revenue by about \$25 billion of gross upside more than offsetting the \$7 billion of potential revenue loss, with the biggest opportunities from money market funds (MMFs), crypto post-trade and derivatives, syndicated loans, private market funds, increased velocity in equities/fixed income and securitized products.

The speed and breadth of adoption will be defined by the automation, speed, liquidity, capital efficiency, and access benefits delivered via tokenization within an FMI perimeter. In addition, regional and regulatory competition is driving adoption in a push to protect or repatriate established revenues pools. The priority for FMIs is to defend post-trade control points in issuance,

registry, clearing, custody, settlement, and servicing, while moving early into asset pools where tokenization solves real friction and building secondary liquidity, new ecosystems, and client connectivity around those.

Prediction markets and Compute contracts point toward the next wave of innovation for derivative franchises, with clear early mover advantages.

Institutional adoption of FMI-addressable, event-risk contracts could become a \$3 billion to \$5 billion annual revenue pool in our base case. Between 60% and 70% of revenue will be net new, coming from financial speculation and hedging as well as deeper penetration of retail and crypto native investors. The remainder will look and feel like traditionally listed derivative contract evolution across macro, policy, weather, corporate and regulated asset-market events and will compete with established listed derivative revenue pools. The revenue opportunity broadens as institutional adoption grows. In our base case, 50% of prediction market-related revenue sits outside of trading, in post-trade, data, and technology franchises.

FMIs need to bolster their market-building muscle and compress product launch cycles.

While product build costs are modest, driving behavioral change, liquidity formation, and scale require different market-building capabilities. This involves distribution, incentives, community engagement, and data advantages, which are areas where crypto and DeFi-native players have so far been more effective than incumbents.

The same priority extends to contracts in new asset classes like Compute where incumbent FMIs can apply the lessons learned from the commodities financialization journey and build Compute into an over-the-counter (OTC) or exchange-traded derivative (ETD) asset class, which could become a \$0.6-1.0 billion revenue opportunity for the FMI sector.

On balance, our base outlook for the sector is positive, but with greater distance between winners and laggards. The sector's "beta" remains attractive with strong baseline growth: Infrastructure revenue continues to benefit from secular deepening across most asset classes, rising complexity of risk management, continued electronification, and even broader retail participation. In some cases, these secular trends may be more powerful than those that propelled FMIs during the last cycle, especially for commodities and listed derivative franchises.

At the same time, the sources of "alpha" are shifting. AI will challenge parts of the data and technology portfolio, but will also create upside from new products, faster product velocity, and cost transformation. Tokenization and prediction markets add further upside by expanding the asset and event universe that can be brought into trusted market infrastructure.

It is back to the future for FMIs in the next cycle with a renewed focus on infrastructure businesses.

The result is a constructive growth outlook in which baseline sector growth remains resilient, but outperformance increasingly depends on business mix, defensibility, and execution. Our base scenario assumes around 8% compound annual growth through 2030. Share of core infrastructure businesses in the overall revenue pool will grow by two percentage points.

In a headwind scenario, sharper AI disruption in commoditized data and workflows, slower new-market adoption, and more limited AI monetization would reduce growth to closer to 6%. In a new-markets upside scenario, growth could be lifted to as much as 12% by faster scaling of tokenized assets and event-risk contracts, as well as other types of new derivative contracts like

Compute, combined with stronger AI monetization in defensible data and workflow products. This could push sector revenue above \$200 billion by 2030, with core infrastructure businesses winning back six percentage points in the sector's business mix.

The winners in the next cycle will be those that are clear about how they organize themselves around four strategic positioning choices in the future FMI ecosystem:

- **Market-of-record.** Winners will prioritize investments into asset classes with the highest secular growth. They will defend trusted post-trade economics as activity moves to new rails and new jurisdictions and build the institutional rails for new tokenized asset classes, event-risk markets, and other new contracts like Compute, before liquidity and data advantages lock in.
- **System-of-work.** Winners will defend embedded technology positions where they sit deep in client workflows, fortifying the embedding and proprietary data and repricing for value. They will build token-aware and event-risk workflow modules to capture new markets.
- **Source-of-origin.** Winners will defend data positions where content is proprietary, authoritative, regulatory-mandated, or embedded into client distribution. They will link exposed analytics into solutions anchored on these moats and reprice toward outcomes. Finally, they will build trusted reference, pricing, and event-data layers for tokenized asset classes and event-risk markets.
- **Operating model.** Winners will treat AI as an operating-model lever, industrializing an "AI factory" that converts cost reduction into EBITDA margin uplift and unlocks the AI revenue upside at scale. They will also close the market-building capability gap (organically or inorganically) relative to crypto and DeFi-native platforms and build the distribution, incentive, community, and product-launch capabilities that consistently turn new asset-class contracts into liquid markets.

01


Industry outlook

This section examines the recent growth of the FMI sector, the structural shifts that will define the next cycle, and our outlook scenarios through 2030. It forms the baseline for the deep dives on artificial intelligence and new markets that follow.

Industrywide revenue has grown by about 15% annually over the last five years, reaching roughly \$131 billion at the end of 2025. This growth was driven by triple tailwinds: First, infrastructure franchises benefited from heightened volatility and secular growth across all asset classes.

Additionally, business mix diversification and increased M&A leveraging the high earnings power of the infrastructure core, expanded data and technology revenues (Exhibit 1). Over the period, the sector shifted its business mix increasingly toward recurring revenues from these new businesses.


Exhibit 1: FMI revenue pool evolution by business line
In \$ billion, 2020–2025



FMI sector revenue grew at a 15% CAGR from about \$66 billion in 2020 to roughly \$131 billion in 2025, led by technology (18% CAGR) and data (25%), while infrastructure (12%) remained the largest pool.
 1.NII, other business lines, e.g. payments.
 Source: Oliver Wyman analysis

Exhibit 2: FMI business mix by archetype

% of total revenue, rounded




■ Infrastructure ■ Technology ■ Data ■ Other

Global FMIs diversified fastest: infrastructure share in revenue mix fell from 75% to 59% of revenue as technology rose from 10% to 22% (14% CAGR, 35% of industry revenue). Regional FMIs stayed infrastructure-led at ~78% (10% CAGR, 53% share). Specialist FMIs grew fastest and most concentrated, infrastructure rising from 91% to 94% (18% CAGR, 12% share).

Source: Oliver Wyman analysis

Exhibit 3: FMI sector revenue growth decomposition by business line

In \$ billion, 2020–2025



Of the roughly \$65 billion of revenue added between 2020 and 2025, infrastructure contributed about \$38 billion, data about \$17 billion, and technology about \$8 billion.


Source: Oliver Wyman analysis

The FMI sector entered this new cycle from a position of strength, with three structural shifts defining the last cycle:

- Volatility and secular capital markets growth.**
 The defining feature of the post-pandemic cycle was an asset-appreciation and volatility-fueled spike in activity across all asset classes, but especially in listed derivatives, commodities, and digital assets. The volatility-driven story has masked strong underlying secular growth drivers, including risk and collateral management, retail participation, capital markets deepening as well as financialization and electrification of asset classes (Exhibit 5). Infrastructure revenue lines added about \$38 billion over the last cycle, growing at around a 12% CAGR.
- Business mix diversification.** Data and technology business lines added \$25 billion of revenue together, growing at 25% and 18% CAGRs respectively (Exhibit 3). 30% of industry revenue is now generated outside traditional infrastructure businesses, driven by expansion into data (including indices, reference data, real-time data, and analytics) and technology (including connectivity, order and execution management systems, trade risk and portfolio management, and RegTech), compared with 21% five years ago. This shift has increased the share of the sector's recurring revenue base and broadened client reach beyond core members to buy-side and corporates.
- M&A.** The sector deployed around \$100 billion of capital, \$88 billion of which was spent on data and technology acquisitions. These were added to boost capability and increase recurring revenue growth (Exhibit 4).

Exhibit 4: FMI M&A deployment by business line

Cumulative, in \$ billion, 2020–2025



FMIs deployed about \$103 billion on M&A between 2020 and 2025 — roughly \$49 billion in technology, \$39 billion in data, and \$15 billion in infrastructure, with around 85% of capital directed at data and technology.

Sources: Oliver Wyman analysis

However, the same forces that enabled this expansive diversification are now evolving. Enterprise AI is stress-testing data and technology acquisitions to determine which are truly synergistic. For instance, connected adjacencies with proprietary content, embedded workflows, regulatory relevance, or infrastructure connectivity should remain defensible, while less connected data, analytics, and shallow workflow tools face greater scrutiny.

AI also offers the sector revenue and operating model opportunities through new products, outcome-based pricing, higher usage, and productivity gains from an industrialized AI factory.

Exhibit 5: Secular growth drivers supporting infrastructure business lines over last and next cycle

Asset class	Growth L5Y	Growth N5Y	Secular growth drivers				Examples
			Risk management	Financialization/ electronification	Retail	Deepening	
Equities	●○○	●●○			✓	✓	• Retail and 24-hour trading • AI and mega-tech wealth creation
Fixed income	●○○	●●○		✓		✓	• Sustained deficits and refinancing wall • Electronification of illiquid credit • ETF and systematic trading
Listed derivatives	●●○	●●●	✓	✓	✓	✓	• Product innovation in event and compute markets • Retail and 24-hour trading • EM market deepening
OTC derivatives	●●○	●○○	✓			✓	• Continued rate-cycle volatility and hedging • Automation and capital efficiency gains
FX	●●○	●●○		✓	✓	✓	• Continued volatility and hedge deepening • FX-as-a-service • EM internationalization
Commodities	●●○	●●●	✓	✓		✓	• Continued financialization and electronification of physical commodities • Energy markets supercharged by AI and green transition
Digital assets	●●●	●●●	✓	✓	✓	✓	• Tokenized RWA starting to scale • Institutional crypto adoption

Growth: ●○○ Low ●●○ Medium ●●● High

Across equities, fixed income, listed and OTC derivatives, FX, commodities, and digital assets, secular drivers, risk management, financialization and electronification, retail participation, and market deepening — continue to support infrastructure growth into the next cycle.

Source: Oliver Wyman analysis

Prediction markets, event-risk contracts, and new contracts for asset classes like Compute then point to the next wave of listed-derivatives innovation, with upside extending beyond trading into post-trade and data and technology services. Three more structural shifts will define the next cycle:

- **AI defensibility and value uplift.** Enterprise AI increases substitutability in commoditized data analytics and non-embedded workflow tools, creating pricing pressure and a higher risk of client insourcing. What remains defensible are revenue pools anchored in proprietary content, authoritative records, regulatory or audit relevance, mission critical workflows, deep client embedding, network effects or trusted infrastructure connectivity. This includes core infrastructure-linked data, indices and benchmarks, regulatory and audit datasets, trade and post-trade workflows, connectivity, risk management, and portfolio-management tools. The exposed pools are those that lack these moats: generic analytics, shallow reporting tools, aggregated historic market data, non-proprietary datasets, and modular workflow components that clients can replicate, substitute, or rebuild using enterprise AI. We estimate 9% of FMI revenue is at direct risk (about \$12 billion), with a further 16% at risk without active mitigation (\$21 billion). At the same time, AI can create \$9 billion of new incremental revenue, equivalent to about 7% of 2025 sector revenue. This is achieved through three levers: \$5 billion from new AI-enabled products, \$2 billion from pricing uplift tied to measurable outcomes, and \$2 billion from higher usage as copilots, agents, and workflow automation increase consumption and client embedding.
 - **AI-enabled operating model.** Enterprise AI also expands the sector's structural margin opportunity. The largest efficiency pools are likely to sit in software engineering and QA, operations, risk, compliance, surveillance, and corporate functions. We see more than 20% cost reduction potential, translating into
- **Tokenization.** Tokenization will reconfigure parts of the existing FMI stack, while expanding the addressable market. Around 35% of today's revenues (\$45 billion) could migrate to new rails, and 5% (\$7 billion) could be structurally eliminated altogether, with the downside most acute where post-trade franchises lose market-of-record status in issuance, registry, clearing, custody, settlement or servicing. Yet broader asset-universe expansion could provide \$25 billion upside, under full penetration, more than offsetting the potential revenue loss. Beyond the well-rehearsed efficiency story, an increasingly relevant adoption driver is regional and regulatory competition, as jurisdictions seek to protect or repatriate clearing and settlement-linked revenue pools. For FMIs, the strategic imperative must be both defensive and offensive, aiming to protect trade and post-trade control points on new rails, while moving early into tokenized asset pools. Here, tokenization will need to solve real frictions around automation, speed, liquidity, and capital efficiency and access, if it is to build new ecosystems and liquidity early.

Taken together, these structural shifts support a healthy growth outlook for the sector but suggest there will be greater distance between winners and laggards.

Exhibit 6: FMI sector outlook scenarios

2025–2030F

Scenario	Features	Outcomes
Base case (Steady infrastructure growth, slow-down in data and technology)	• Core infrastructure grows steadily, supported by continued secular capital markets growth vectors • Data and technology growth moderates from prior highs due to pricing pressure, maturing penetration in core segments, and only limited AI-driven upside • Revenue mix remains broadly stable, with infrastructure growth offsetting slower expansion in data and technology	• Sector growth remains stable at a high single-digit rate overall • Profitability is supported by operating discipline, pricing management, and incremental innovation within existing models • Strategic posture stays focused on optimization of the current portfolio rather than major business model change
AI headwind (Sharper disruption in data and technology)	• Data and technology revenues face pressure from AI-driven compression in both pricing and usage, especially in less differentiated products • Core infrastructure remains comparatively resilient due to network effects, trust, embedded workflows, yet tokenization slowly eating into infrastructure revenues • Investment priorities shift toward more differentiated, defensible, and higher-value data and technology offerings as well as defense of infrastructure revenues	• Sector growth becomes more polarized, with steady infrastructure performance but weaker data and technology growth • Profitability benefits from enterprise AI adoption and internal efficiency gains, despite external revenue pressure in exposed products • Strategic posture shifts toward repositioning portfolios around defensible capabilities and away from commoditizing offerings
New markets scale (New market starting to scale alongside AI monetization)	• Infrastructure revenue pools are retained and significantly expanded via tokenization and scaling of prediction markets • Data and technology revenue pressure remains fierce but is slightly moderated by AI-related product and pricing upside getting realized, and data and technology spill-overs from new markets • Revenue mix shifts significantly towards infrastructure related revenues on the back of a significantly expanded asset universe and new markets	• Sector growth accelerates, with potential for double-digit infrastructure growth and partially mitigated AI headwinds • Profitability improves through operating leverage from new revenue streams, higher infrastructure growth, and scaling benefits from tokenization and adjacent market expansion • Strategic posture shifts from defending the core to scaling new businesses and capturing adjacencies

Three scenarios frame the 2025–2030 outlook: a base case of steady infrastructure growth with slowing data and technology; an AI-headwind case of sharper data and technology disruption; and a new-markets upside case in which tokenization and new markets scale.

Source: Oliver Wyman analysis


The sector's beta also remains highly attractive. While players can expect some continued volatility, there are significant secular trends supporting the continued growth of infrastructure businesses. The need for risk management, ongoing financialization and electrification of asset classes, deeper retail penetration, and general capital market deepening will continue and, in some cases, even accelerate. The asset classes with the highest relative growth will be listed derivatives, commodities, and digital assets (Exhibit 5).

Alpha will depend on two execution agendas. The first is AI: defending exposed data and workflow revenues, monetizing AI through new products, pricing, and usage, and converting productivity gains into margin expansion. The second is new markets: protecting market-of-record positions as tokenization migrates parts of today's post-trade activity to new rails, while capturing upside from tokenized asset pools, event risk, and other contracts. We have created three scenarios for how we expect the industry to fare between now and through 2030 (Exhibit 6).

- Scenario A “Base case”:** This scenario anticipates secular infrastructure growth, manageable AI disruption in contestable data and shallow workflow pools, gradual AI monetization and cost transformation, and measured adoption of tokenization and prediction markets. Result: around 8% annual growth through 2030, with outperformance increasingly dependent on business mix, defensibility, and execution. Share of core infrastructure businesses in the overall revenue pool will grow by two percentage points.
- Scenario B “AI headwind”:** This scenario is based on sharper AI disruption in commoditized data and shallow workflows, slower adoption of tokenized assets and event-risk contracts, and more limited AI monetization. Core infrastructure resilience limits the downside, with moderate sector growth at around 6% annually through 2030.
- Scenario C “New markets scale”:** This scenario depends on faster scaling of tokenized assets and event-risk contracts, stronger AI monetization in defensible data and workflow products, and more effective AI-led cost transformation. In this scenario, the sector could grow as much as 12% annually, and sector revenue could expand to more than \$200 billion by 2030, with core infrastructure businesses winning back six percentage points in the sector’s business mix.

Exhibit 7: FMI revenue outlook by scenario

In \$ billion, 2025–2030F



Sector revenue could grow from \$131 billion in 2025 to about \$196 billion in the base case ($\approx 8\%$ CAGR), \$177 billion under an AI headwind ($\approx 6\%$), or \$229 billion in a new-markets upside ($\approx 12\%$) by 2030.

Source: Oliver Wyman analysis

Artificial intelligence

DEEP DIVE

This section explores how enterprise AI is reshaping FMI business models and performance. Across infrastructure, data, and technology, we assess where AI poses a disruption risk, durable moats protect value, and AI can create revenue upside through new products, pricing uplift, and higher usage. We then estimate the size of the opportunity created by AI-driven operating leverage and provide perspectives on what FMI management teams need to do to scale AI enterprise-wide, including ideas on approaches to risk management and governance.

2A. AI defensibility

Enterprise AI changes the outlook for data and technology business lines, increasing substitutability and putting pressure on pricing where products lack durable moats.

For the FMI sector, the key questions are which revenue pools are protected by proprietary content, embedded workflows, network effects, regulatory relevance, or source-of-origin status, and which pools are exposed to faster commoditization and client insourcing.

The answers directly affect the diversification thesis of the last cycle. Strategically and operationally synergistic adjacencies should prove valuable and defensible, while less connected extensions, acquired during the last cycle that are disproportionately concentrated in commoditized data, analytics, and workflow risk becoming challenged, or a drag on FMI group value.

Exhibit 8 provides a defensibility lens across FMI business lines, structured around six core dimensions that determine how resilient a revenue pool is to AI-driven disruption. This looks at:

- **Network effects.** Value increases as more participants, counterparties, and data contributors join, reinforcing liquidity, usage, and data advantages including exchange liquidity pools, and CCP clearing networks.
- **Workflow embeddedness.** Products are tightly integrated into client processes, decision loops, and operational workflows. This increases stickiness through order management systems and post-trade processing workflows.

- **Mission criticality.** Solutions are essential to core operations, where failure or downtime carries high financial, operational, or regulatory risk.
- **Regulatory and audit roles.** Products provide compliance, auditability, and control and reporting capabilities required for market participation, including trade reporting, surveillance, and regulatory data services.
- **Deep proprietary IP.** Value is anchored in unique data, models, or content that cannot be easily replicated using general-purpose AI, such as used in index methodologies, high-quality reference data, and pricing curves.
- **Trust and authority.** Institutions rely on the provider as a source-of-origin or neutral utility, with strong brand, credibility, and reliability. This comes to play in official price formation, benchmark indices, and central securities depositories.

Based on this assessment, we classify FMI business lines into three categories with respective management actions (Exhibit 9):

- **Fortress (75% of FMI revenue, \$98 billion). Defend and price for mission-critical outcomes.** Pools with the strongest defensibility across the six dimensions, where value is anchored in trust, regulation, and mission-critical functionality, enabling pricing based on outcomes and reliability. Examples include issuance, listing, trading, clearing, and settlement, indices and benchmarks, connectivity, trade risk management, and portfolio management.

Exhibit 8: AI defensibility framework for FMI revenue pools

Business line		Network effects	Workflow embeddedness	Mission criticality	Regulatory and audit role	Deep proprietary IP	Trust and authority	Total
Infrastructure	Issuance							
	Listing							
	Trading							
	Clearing and risk management							
	Settlement and custody							
Data	Real-time market data							
	Historical market data							
	Reference data							
	Indices and benchmarks							
	Analytics							
	Other data							
Technology	Connectivity							
	Order and execution management							
	Trade risk management							
	Portfolio management							
	Analytics and reporting							
	Data management							
	Regulatory workflow							

AI defensibility



FMI business lines are scored from low to very high across six dimensions of AI defensibility: network effects, workflow embeddedness, mission-criticality, regulatory and audit role, deep proprietary IP, and trust and authority.

Source: Oliver Wyman analysis

- **Augment (16% of FMI revenue, \$21 billion).**
Embed AI to defend value, which is otherwise at risk without active mitigation. Pools that remain structurally defensible but require AI to enhance value, improve efficiency, and sustain differentiation. Examples include real-time market data, reference data, order and execution management, and regulatory workflow. These are the pools where outcome-based repricing and embedded agents are most important.
- **Contest (9% of FMI revenue, \$12 billion).**
Redesign, bundle, or exit because it is at direct risk. Pools with weaker defensibility in which AI increases substitutability and pricing pressure, requiring redesign, bundling, or selective exit. Examples include non-differentiated data analytics, analytical and reporting software (non-embedded), data management (non-curated), historic market data, and non-proprietary data.

Exhibit 9: FMI revenue exposure to AI disruption

% of 2025 revenue

Summary by management action

 Fortress Defend and price for mission critical outcomes	 Augment Embed AI to improve client outcomes	 Contest Re-design, link or exit
• Issuance, listing, trading, clearing and settlement • Indices and benchmarks • Connectivity • Trade risk management • Portfolio management (including xBOR) • ...	• Real time market data • Reference data • Order and execution management • Regulatory workflow • ...	• Data analytics • Analytical/reporting software • (Non-embedded) data management • (Non-curated) historic market data • Other (non-proprietary) data • ...
75% of industry revenue	16% of industry revenue	9% of industry revenue

Applying the defensibility lens, about 75% of FMI revenue is Fortress (defend and price for outcomes), about 16% Augment (embed AI), and about 9% Contest (redesign, bundle, or exit).

Source: Oliver Wyman analysis

Refreshing the M&A agenda for the next cycle

Enterprise AI should prompt FMI groups to revisit the M&A playbook of the last cycle. A meaningful share of the c. \$100 billion deployed between 2020 and 2025 went toward data analytics, reporting software, data management, and aggregated datasets, reflecting a belief that scale in data would provide durable pricing power and differentiation. AI is now testing that thesis, increasing the value of proprietary content, embedded workflows, regulatory relevance, trusted infrastructure positions, and network effects, while reducing the defensibility of more generic data, analytics, and reporting capabilities.

The key lesson is that assets need to be both defensible and strategically or operationally synergistic with the core FMI franchise. Acquisitions that strengthen a firm's position as a market-of-record, system-of-work, or source-of-origin should become more valuable because they reinforce existing moats, deepen customer integration, and create opportunities for cross-sell and product innovation. By contrast, businesses built around aggregated, non-proprietary content, generic analytics, or shallow workflow functionality will face increasing pressure, particularly when they lack clear links to core infrastructure, proprietary data, or critical client workflows.

As a result, the M&A agenda should shift from acquiring scale to acquiring defensibility and synergy. Management teams should reassess existing portfolios to determine which acquired businesses can be strengthened through proprietary data, workflow integration, and AI-enabled innovation, and which should be bundled, divested, or deprioritized. Future acquisitions should focus on capabilities that reinforce durable moats, strengthen strategic control points, and expand relevance in emerging markets, such as tokenized assets and event-risk infrastructure.

In the AI era, successful M&A will be defined less by diversification and more by the combination of defensibility, connectivity to the core franchise, and a differentiation of platforms that AI cannot easily replicate.

Actions for FMI management

- Segment data pools by defensibility
- Protect and reprice Fortress data
- Redesign, link, or exit exposed pools

2B. AI value uplift

Enterprise AI also creates a meaningful revenue upside for FMIs. The opportunity is not uniform and concentrates in pools with strong defensibility (Fortress and Augment) in which products can embed AI into workflows and outcomes. We see three primary levers of upside: new products, pricing, and usage, which can often have a compounding impact when executed together (Exhibit 10).

At a sector level, our base view (Exhibit 11) implies a revenue uplift opportunity of around \$9 billion from AI-specific impacts, including an estimated \$5 billion from new products, \$2 billion from pricing, and \$2 billion from usage.

- **New products (+\$5 billion, +4%):** AI enables the creation of differentiated, higher-value offerings built on existing data, infrastructure, and workflow positions. This includes AI-enhanced analytics and signals on top of real-time data, surveillance-as-a-service, automated post-trade exception handling, and decision-support tools embedded in order management systems/execution management systems (OMS/EMS).
- **Pricing (+\$2 billion, +1–2%):** AI shifts pricing power toward outcomes, such as pricing based on risk reduction, time saved, or compliance readiness, rather than seats or feeds. Providers with strong regulatory and audit roles and mission-critical workflows can re-anchor pricing, introduce outcome-based tiers, and bundle AI capabilities into premium packages. For the most differentiated data businesses, the re-pricing opportunity will be more than 30% like-for-like.
- **Usage (+\$2 billion, +1–2%):** AI increases throughput and frequency of use, including through more queries, automated workflows, and broader user personas, expanding consumption-based revenues with application programming interfaces (APIs), analytics calls, and data usage, and driving higher wallet share within existing clients. Embedded copilots and agents can materially increase daily active usage of data and workflow products.

The largest AI revenue upside sits where FMIs can combine trusted market infrastructure, proprietary data, and embedded client workflows into higher-value products. First, new AI-enabled products can be built on top of existing control points, such as AI-enhanced analytics and signals, surveillance-as-a-service, automated post-trade exception handling, and decision-support tools embedded in OMS/EMS workflows. Second, providers can reprice around measurable outcomes where they reduce risk, save time, improve compliance readiness, or automate mission-critical processes. Third, copilots and agents can increase usage by expanding the frequency, depth, and number of users interacting with data, analytics, and workflow products.

The biggest opportunities are not in standalone AI features, but in AI embedded into products that already benefit from trust, authority, regulatory relevance, proprietary content, or workflow integration.

Exhibit 10: AI revenue opportunity by lever

2025–2030F

Business line		New products	Pricing	Usage	Total
Infrastructure	Issuance				
	Listing				
	Trading				
	Clearing and risk management				
	Settlement and custody				
Data	Real-time market data				
	Historical market data				
	Reference data				
	Indices and benchmarks				
	Analytics				
	Other data				
Technology	Connectivity				
	Order and execution management				
	Trade risk management				
	Portfolio management				
	Analytics and reporting				
	Data management				
	Regulatory workflow				


Low Very high

AI revenue upside, scored by business line across three levers: new products, pricing, and usage and concentrates in analytics, regulatory workflow, portfolio management, reporting, and reference and historical data, and is smallest in connectivity, adding about \$9 billion in aggregate.

Source: Oliver Wyman analysis

Exhibit 11: AI revenue upside by lever

In \$ billion, 2025



Enterprise AI could lift sector revenue from \$131 billion to about \$140 billion, roughly \$9 billion of upside: through new products, pricing, and higher usage.

Source: Oliver Wyman analysis

Actions for FMI management

- Embed AI into defensible workflow pools
- Reprice AI-enabled offerings toward outcomes
- Drive usage uplift through copilots, agents, and automation

2C. AI-enabled operating model

In parallel to revenue-side impacts, AI represents a material opportunity for FMIs to optimize operating models. The sector has a structural advantage in capturing productivity uplift, given its bias toward standardization and scale. However, a meaningful share of activity sits in regulated or supervisory-critical functions, where model risk, auditability, explainability and resilience requirements raise the bar for AI deployment. As a result, near-term adoption is typically faster in corporate and enablement functions before scaling into market-facing functions.

Exhibit 12 illustrates the AI operating model opportunity across FMI functions. We see more than a 20% cost reduction potential, driven by a combination of automation and operating leverage. The largest opportunities are concentrated in software engineering and quality assurance (where a significant portion of the FMI cost base sits), operations, risk, compliance, surveillance, and corporate functions. Early value is driven by assistive use cases that increase throughput, while the medium-term prize comes from end-to-end process and ultimately organizational redesign.

Across functions, uplift is driven by four core mechanisms:

- **Cycle-time compression:** fewer handoffs and faster resolution of breaks and incidents such as with agentic onboarding and know your customer (KYC) orchestration, real-time client servicing, and faster data reconciliation.
- **Higher first-time-right rates:** better triage, stronger data validation, and fewer downstream exceptions in such areas as automated QA and reconciliation, fraud detection, and dynamic client risk assessment.
- **Lower unit cost-to-serve:** automation of repeatable tasks and knowledge work such as with AI-assisted software development, continuous forecasting, AI-enabled third-party management, and client communications.

- **Better reuse at scale:** common components like models, prompts, and retrieval layers deployed across products and jurisdictions in shared models for sentiment analysis, surveillance, and client touchpoints across front office and business lines, risk management and marketing.

Capturing this opportunity is not a tooling exercise. In FMI environments, value only scales when AI is industrialized through reusable components, strong governance, and clear productivity KPIs, effectively an “AI factory.”

This implies standardized model pipelines, shared data and feature layers, robust model risk management, and a control tower to prioritize use cases, track value, and enforced reuse across functions. Without this factory approach, deployments are likely to remain fragmented pilots, with duplicated spend across complex organizations.


Exhibit 13 translates these efficiency gains into an operating leverage pathway. Sector EBITDA margins have been broadly stable at 52–54% over the last cycle, but AI can enable a step-change to more than 60% by 2030, implying a nine-percentage-point increase. This reflects both direct cost savings (driven by a more than 20% cost reduction opportunity) and additional operating leverage from more scalable, AI-enabled operating models.

Exhibit 12: AI cost opportunity by FMI function

% cost reduction potential by 2030F

Cost profile development of an average FMI¹

2025–2030E, in %



Selected AI efficiency drivers by function

Illustrative, non-exhaustive

Function	
Front office ²	• Agentic onboarding (e.g., agent orchestrated KYC as a transversal capability across businesses) • AI-enabled client touchpoints and augmented sales force (e.g., self-serve portals, next-best action engines embedded across client-facing roles) • Augmented client servicing (e.g., real-time call transcript analytics, pro-active issue resolution)
IT	• AI-enabled software development (e.g., Copilot for developers) • Agentic acceleration of data reconciliation and quality assurance • Predictive maintenance of hardware
Finance	• Continuous forecasting via real-time dashboards • AI-enabled 3rd-party management (e.g., streamlining of KYS processes and existing networks)
HR	• AI-enabled capability management (e.g., personalized, assisted onboarding, and learning journey, self-serve procedures) • AI-supported performance management (e.g., review writing)
Legal, risk and compliance	• AI-enabled transaction monitoring and fraud detection • Dynamic client risk assessment (e.g., periodic KYC review)
Marketing and comms	• AI-powered issuer and client communication (e.g., market updates, regulatory updates) • AI-driven sentiment and general client behavior analysis

AI could cut an average FMI's cost base by about 20% to 25% by 2030, with reductions of roughly 20-30% across front office, IT, finance, HR, legal, risk and compliance, and marketing.


1. Assumption: Integrated FMI with combined infrastructure, data, and technology business; 2. Including sales, customer service, and product development.

Source: Oliver Wyman analysis

Capturing this opportunity is the foundational element of the operating model agenda for FMIs in the next cycle: It converts cost reduction into margin expansion through the AI factory, and unlocks AI revenue upside in workflow products at scale. This requires an AI factory that can convert cost reduction into margin expansion.

The model also unlocks the AI revenue upside in workflow products at scale. To be successful, it requires a fundamental rethink of how FMIs are organizationally structured, partially moving away from organizational silos and tailored ways-of-working that deep specialization often necessitated.

Exhibit 13: AI-enabled EBITDA margin uplift potential
 % EBITDA margin, 2025–2030F



AI-driven efficiency could add about nine percentage points to FMI EBITDA margins by 2030, on margins that range by archetype from roughly 36% at specialist FMIs to about 61% at global FMIs.
 Source: Oliver Wyman analysis

Actions for FMI management

- Create a top-down cross-organizational vision for AI
- Industrialize AI through a factory model
- Convert AI productivity into margin uplift

2D. AI governance and risk

The ambition to scale AI across financial institutions is moving from narrow productivity tools toward larger volumes of use cases and more complex agentic or autonomous applications. For FMIs, the opportunity is material, but the adoption path is structurally more constrained than in many other parts of financial services. FMIs operate under standards and supervisory expectations that place the highest expectations on governance, operational resilience, cyber recovery, settlement continuity, and control of critical third-party dependencies. That raises the bar for using AI in core decision workflows and explains why adoption remains comparatively cautious.

Where FMIs are using AI today, adoption is still concentrated in internal support activities, with limited autonomy and relatively rare use in core market, risk, or operational decision processes. The constraint is not only technology readiness, it is whether FMIs can build risk and governance frameworks that are reliable enough to support scaled deployment without compromising resilience, auditability, or trust.


Traditional governance models are not sufficient for this next phase. AI systems can be developed quickly, deployed across many use cases and, increasingly, operate with a degree of autonomy. Static approvals, periodic reviews, and manual policy policing will become a bottleneck if AI use cases proliferate across the enterprise. Thus, FMIs need to shift from episodic model review toward continuous, embedded governance and controls, supported by automation-first assurance, real-time monitoring, and clear ownership across the lines of defense.

The three lines-of-defense operating model needs to change in three ways. First, FMIs need a top-of-house AI risk and governance framework that defines risk appetite, use-case classification, approval paths, controls, monitoring standards, and escalation routes across all AI applications. Second, the lines of defense need clearer mandates. The first line should be accountable for AI solutions and their business outcomes, not just ad hoc model ownership, while the second line should move from manual challenges to framework stewardship, orchestration, and targeted assurance. Third, testing, validation, and monitoring need to become scalable and continuous, covering transversal risks such as model drift, explainability, data leakage, cyber vulnerability, third-party dependency, operational resilience, and unintended autonomous behavior (Exhibit 14).

**This governance approach
should be an enabler of AI
adoption, not a constraint on it.**

Without it, FMIs will struggle to move beyond low-risk internal support tools and will be slow to capture the revenue and margin opportunity set out above. With it, they can scale AI into more valuable workflow, risk, compliance, surveillance, and operations use cases while preserving the trust, continuity and control standards that define the sector.

Exhibit 14: AI risk and governance framework



As FMIs move from single-task models to integrated, autonomous AI systems, risk management must shift from use-case checks to system-level testing, validation, and continuous runtime monitoring across transversal risks.
Source: Oliver Wyman analysis

Actions for FMI management

- Make AI risk and governance scalable

03

New markets

DEEP DIVE

This section explores how new markets are reshaping the FMI perimeter. It examines the impact of tokenization on today's infrastructure stack, the asset pools where tokenization could expand FMI relevance, the potential for prediction markets to evolve into institutional event-risk infrastructure, and lastly the potential of Compute markets.

3A. Tokenization

For FMIs, tokenization creates both upside and downside.

It could shift issuance, settlement, custody, and servicing revenues from legacy post-trade processes to token-native rails and trading as well as associated data and technology revenues to digital venues. The tokenization opportunity is best viewed through a perimeter lens: which asset classes already sit inside the FMI ecosystem, and which large pools remain outside because they are fragmented, bespoke, illiquid, or operationally complex.

Tokenization can expand the addressable market by making more assets easier to issue, transfer, collateralize, service and, in some cases, trade. We expect the expansion effect to outweigh displacement, but only for FMIs that move early enough to become trusted infrastructure providers for the new rails and build ecosystems and connectivity to new clients.

If incumbents do not establish themselves on the new rails, setting the standards, holding the trusted records, and operating the core data layers, a meaningful share of today's post-trade economics will migrate to token-native infrastructure providers.

Exhibit 15 frames the impact on the existing FMI revenue pool across three categories:

- Revenue untouched (about 60% of FMI revenue, \$79 billion):** A large share of existing revenue remains structurally resilient because tokenization does not remove the need for trusted markets, liquidity formation, risk management, data analytics, or workflow tools. Core venue revenues, listings, risk-management revenues, prices, histories, benchmarks, analytics, connectivity, OMS/EMS, and portfolio tools continue to be required, even if the asset format changes.

Exhibit 15: FMI revenue exposure to tokenization
% of 2025 revenue

Revenue untouched	Revenue migrated to new rails	Revenue structurally eliminated
Core venue, listing and risk-management revenues remain essential even in tokenized markets	Issuance, trading, clearing, settlement and custody economics shift onto token-native rails and platforms	Only a small slice (mainly NII and legacy friction fees) is structurally removed
Most spend on connectivity, risk, OMS/EMS and portfolio tools continues in tokenized markets	A material share shifts into token-aware modules, integration layers and on-chain workflow tools	Only niche legacy systems vanish where smart contracts remove entire process steps
Demand for prices, histories, benchmarks and analytics largely persists regardless of asset format	Some economics move to providers of on-chain price, reference and token-specific data	Very little data revenue disappears; tokenization changes sources more than overall demand
60% of FMI revenue	35% of FMI revenue	5% of FMI revenue

■ Infrastructure ■ Technology ■ Data

Tokenization leaves about 60% of FMI revenue untouched, shifts about 35% onto new token-native rails, and structurally eliminates only about 5%, with the impact concentrated in post-trade infrastructure economics.

Source: Oliver Wyman analysis

- **Revenue migrated to new rails (about 35% of FMI revenue, \$45 billion):** A meaningful share of economics may shift to new digital rails or venues. Issuance, clearing, settlement and custody revenues as well as select trading economics may migrate toward token-native rails and platforms; data economics shift toward providers of on-chain price, reference and token-specific data; and technology spend shifts into token-aware modules, integration layers and on-chain workflow tools.
- **Revenue structurally eliminated (about 5% of FMI revenue, \$7 billion):** Only a small slice of economics is structurally removed, for example where tokenization eliminates legacy frictions, manual process steps or certain NII-like or float-driven economics.

The new driver: regional and regulatory competition

The dominant industry narrative on tokenization is an efficiency story, focused on faster settlement, more automated lifecycle management, lower capital cost, easier distribution efforts, and lower operational cost. Through that lens, the adoption portfolio looks narrow, and pure efficiency gains tend to be competed away as adoption scales.

A new material driver of tokenization's strategic significance is regional and regulatory competition. Multiple jurisdictions are actively positioning their domestic infrastructure stack to repatriate clearing, settlement, and custody revenues currently captured in established end markets, with sovereign incentives that go well beyond efficiency.

For incumbent FMIs, the strategic calculus changes in three ways. First, the question is no longer whether to participate in tokenization, but where to invest to retain a market-of-record role in jurisdictions that matter. Second, the cost of inaction is asymmetric: The loss is not a slow erosion of revenue, but a discrete shift of regulatory perimeter that, once set, is extremely hard to reverse.

Third, partnerships with sovereign and quasi-sovereign actors, including central banks, treasuries, and regulators, become a central part of the strategy. Most exposed are FMIs with post-trade economics that depend on a small number of jurisdictions where regulatory change is moving fastest.

Best positioned are those with multi-jurisdictional footprints and the credibility to be selected as trusted infrastructure providers on the new rails. Yet capability building is not a quick exercise, but could take years.

Set against the defensive picture, tokenization creates significant revenue opportunity by expanding the universe of assets that can be served by institutional market infrastructure.

The opportunity is not to tokenize everything, but to target large, friction-heavy asset pools where today's processes remain manual, fragmented, or hard to scale.

In these markets, tokenization can reduce lifecycle complexity, accelerate settlement, improve secondary liquidity, and capital efficiency, broaden investor access, and create more standardized data and ownership records.

For FMIs, this creates new infrastructure, data, and workflow revenues. The upside is concentrated in asset classes where tokenization solves real market frictions rather than simply creating a new wrapper for already efficient instruments. Exhibit 16 assesses nine asset pools across five benefit dimensions: lifecycle automation, settlement speed, secondary liquidity, capital efficiency, and broadened access.

The largest opportunity pools are those with both large existing balances and multiple tokenization benefits, which will drive near term adoption:

- **The highest-priority opportunities** are money market funds, crypto post-trade and derivatives, syndicated loans, and increased velocity in equities and fixed income. MMFs benefit from tokenized fund shares that can improve collateral mobility, settlement speed, and integration into liquidity-management workflows. Crypto post-trade and derivatives are already token-native or adjacent but remain under-institutionalized.

The opportunity for FMIs is to bring trusted clearing, settlement, custody, margining, surveillance, market data, and risk controls to a market that has scale but lacks institutional plumbing. Syndicated loans remain operationally complex, with delayed settlement, manual transfers, and fragmented servicing, making them a natural candidate for tokenized ownership records, entitlement management, and secondary-market infrastructure. Moreover, tokenization will increase the distribution reach and velocity across equities and fixed income, specifically with retail and hedge funds.

- **A second opportunity tier** consists of private market funds, securitized products, and pre-IPO equity. These pools are sizable and operationally heavy, but adoption is likely to be more gradual and dependent on issuer, investor, and regulatory readiness. Tokenization can still create meaningful FMI-relevant revenues by standardizing issuance, registry, transfer agency, ownership restrictions, servicing, pricing, and lifecycle data. The strategic appeal is that these assets currently sit only partly inside the institutional market-infrastructure perimeter. Tokenization could make them more transferable, data-rich, and compatible with market-style infrastructure over time.
- **A third tier** covers commodities and payments/stablecoin infrastructure. These opportunities are more targeted. In commodities, the value case is strongest where tokenization improves custody, provenance, collateral eligibility, audit trails, or access to fractional ownership. In payments and stablecoin infrastructure, the opportunity is less about tokenizing an asset pool and more about providing trusted settlement, interoperability, compliance, and connectivity layers around tokenized cash, reducing frictions in combination with clearing and settlement processes.

The implication is that FMIs should prioritize tokenization based on revenue relevance and control-point potential, not just theoretical asset balances.

Exhibit 16: Tokenization benefit assessment by asset pool

Asset class	Tokenization benefits					Summary assessment
	Lifecycle automation Smart contract replaces manual events (e.g. capital calls, coupons, amendments)	Settlement speed Near-instant versus multi-day or multi-week today	Secondary liquidity Creates tradability for currently illiquid or thinly-traded assets	Capital efficiency Collateral efficiency and liquidity/funding benefits	Broadening access Lowers minimums or opens institutional access to currently restricted pools	
Money market funds	✓	✓		✓		●●●●
Crypto post-trade and derivatives	✓	✓		✓	✓	●●●●
Increased velocity of existing assets (EQ and FI)	✓	✓		✓	✓	●●●●
Syndicated loans	✓	✓	✓		✓	●●●●
Private markets funds	✓	✓	✓		✓	●●●●
Securitized products	✓		✓	✓	✓	●●●●
Pre-IPO equity	✓	✓	✓		✓	●●●●
Commodities	✓			✓	✓	●●●●
Payments and stablecoins (infrastructure)	✓	✓		✓		●●●●


Driver: ✓ Primary ✓ Secondary **Benefit:** ○○○ Low ●●● High

Nine asset pools are assessed across five tokenization benefits: lifecycle automation, settlement speed, secondary liquidity, capital efficiency, and broadening access, with money market funds, crypto post-trade and derivatives, syndicated loans.

Source: Oliver Wyman analysis

Exhibit 17: Tokenization revenue upside by asset class and value-chain step

In \$ billion, under full penetration



Asset class

■ Private market funds ■ Securitized products ■ Pre-IPO equity ■ Syndicated loans ■ Money market funds
 ■ Payments and stablecoins (infrastructure) ■ Crypto post-trade and derivatives ■ Commodities ■ EQ and FI velocity

Value chain

■ Issuance/registry ■ Listing ■ Trading/transactions ■ Clearing ■ Custody ■ Settlement ■ Collateral management

After about \$7 billion of revenue is structurally eliminated, roughly \$25 billion of tokenization upside grows the FMI revenue pool from \$131 billion to about \$150 billion, concentrated in post-trade steps across crypto, syndicated loans, and money market funds.

Source: Oliver Wyman analysis

At a sector level, our base view is that tokenization can expand the FMI revenue pool from \$131 billion in 2025 to as much as \$150 billion under full penetration. This reflects two offsetting forces: approximately \$7 billion of structurally eliminated revenue, which can be more than offset by approximately \$25 billion of gross upside from new tokenized asset pools and token-native infrastructure services.

The largest overall opportunity is crypto post-trade and derivatives: The market is already scaled and token-native, but still lacks institutional clearing, settlement, custody, margining, surveillance, and data and risk controls. Syndicated loans are the next major pool, with high benefits from digitizing ownership records, entitlements, transfers, servicing, and secondary-market infrastructure.

Money market funds are smaller in overall opportunity size but offer the clearest near-term revenue case because tokenized fund shares can be integrated into collateral, settlement, and liquidity-management workflows quickly. From a value chain perspective, the opportunity is concentrated. Roughly 75% sits in post-trade functions and 25% in trading. The biggest post-trade pools are issuance and custody with clearing, settlement and collateral management following. Trading remains important, but the larger prize is building the institutional post-trade infrastructure that allows tokenized markets to scale.

For incumbent FMIs, this creates a different set of strategic choices than prior technology cycles. The defensive priority is to protect existing trade and post-trade economics as settlement, custody and registry functions move onto new rails and, increasingly, into new jurisdictions. The offensive priority is to expand the infrastructure perimeter into assets that historically sat outside the FMI ecosystem because they were too bespoke, fragmented, or operationally intensive to support at scale.

Actions for FMI management

- Defend control points on new rails and build ecosystems
- Prioritize tokenization in areas in which revenue relevance and benefits are highest
- Treat regional and regulatory competition as a first-order issue

3B. Prediction markets

Prediction markets are still early in their maturity, but they are no longer just a retail curiosity. Volumes today are concentrated in politics, crypto, sports, and entertainment-adjacent use cases, but the underlying market structure has broader relevance. At their simplest, prediction markets turn uncertainty into tradable, observable probabilities. For FMIs, the relevant question is whether enough FMI-relevant events can become standardized, liquid, and well-governed markets.

The near-term opportunity is likely to remain narrow. Many contracts will stay retail-led and concentrated on sports and entertainment. But a smaller set of categories has a clearer fit with FMIs: macro releases, central-bank decisions, inflation outcomes, weather events, corporate actions, initial public offerings,

M&A outcomes, and regulated asset-market events. These are areas where investors, corporates and risk managers already have exposures but often lack a clean instrument to price or hedge a specific outcome.

Exhibit 18 distinguishes between where prediction-market volume sits today and where FMI addressability is highest. Current activity is concentrated in retail-led categories such as politics and elections, crypto and digital assets, sports, entertainment, and culture, which have proven engagement and liquidity formation but are not always the most natural fit for incumbent FMIs. The more FMI-addressable opportunities sit in categories that connect to institutional risk, data, and workflow needs: macro and financial-market events, weather and climate, corporate and real-world events, and regulated asset-market events.

Exhibit 18: Prediction market contract taxonomy and use cases

FMI addressability	Contract coverage		Example contracts	Current volume
	Politics and elections	Elections, party races, leadership contests, policy outcomes, government decisions	"Who wins the election? Will there be a government shutdown?"	
	Sports	Match outcomes, championships, player awards, season performance	"Who wins the NBA Finals? Will Team X make the playoffs?"	
	Entertainment and culture	Awards, box office, TV renewals, celebrity and media-driven events	"Who wins Best Picture? Will Show X be renewed?"	
	Crypto and digital assets	Token prices, ETF approvals, protocol upgrades, airdrops, crypto regulation	"Will BTC exceed 100k? Will an ETH ETF be approved?"	
	Weather and climate	Temperature, rainfall, storms and other measurable climate events	"Will rainfall exceed Y? Will July temperatures be above average?"	
	Corporate and real-world events	M&A approvals, IPOs, launches, legal outcomes and company milestones	"Will Company X complete acquisition Y? Will Startup Z IPO this year?"	
	Macro and financial markets	Rates, inflation, jobs data, recessions, equity indices, FX and commodities	"Will the Fed cut rates? Will CPI print above X?"	

○ No volume ● All volume

Prediction-market contracts span seven categories, from politics and sports to macro and financial markets; today's volume is concentrated in retail categories, while FMI-addressable opportunity is highest in macro, corporate, and weather events.

Source: Oliver Wyman analysis

Four categories may have lower volumes today, but they are more relevant for FMIs because they can be governed with objective settlement criteria and linked to existing trading, hedging, clearing, data, and risk-management workflows:

- **Macro and policy events:** contracts on rate decisions, inflation prints, jobs data, recessions and FX or equity index levels can create market-implied probabilities for events that already drive large hedging and allocation decisions.
- **Weather and climate events:** contracts on temperature, rainfall, storms, and other measurable outcomes can provide a simpler risk-transfer tool for exposures currently managed through insurance, bespoke derivatives, or not hedged at all.
- **Corporate and real-world events:** contracts on M&A approvals, IPO timing, product launches, legal outcomes, and company milestones could support event-driven investors and corporates, provided the contracts are carefully governed and tied to objective settlement criteria.
- **Regulated asset-market events:** contracts linked to crypto ETF approvals, protocol events, index inclusion, benchmark changes, or other market-structure events can create transparent pricing for risks that already trade indirectly through related assets.

The market-building capability gap

The cost of building a prediction-market product is modest. Contract design, settlement logic, surveillance, and connectivity sit well within the existing capability set of any major exchange or clearing house. The hard part, and the part incumbents have systematically under-invested in, is driving the behavioral change required for participants to adopt these contracts in the first place and reach the scale at which network effects, liquidity, and data advantages kick in. That is a different kind of problem, and it requires a different kind of spend and a different kind of thinking.

Crypto and DeFi-native platforms have so far been materially more effective at this than established incumbents. Their playbook combines direct-to-user distribution through web, mobile and social channels rather than reliance on member-firm sales; structured incentive programmes (liquidity mining, market-maker subsidies, token-aligned rewards) that bootstrap two-sided activity; open community engagement, public roadmaps and contract design driven by user demand; and the use of liquid probabilities as a data

and media product, amplified through news cycles in ways that create reflexive volume growth. None of these are technical innovations. They are sustained investments in behavioral change, distribution and community, the things that turn a listed contract into a liquid market.

Incumbents bring strengths crypto-native platforms cannot easily replicate, such as rulebook credibility, clearing infrastructure, surveillance, regulated settlement, and trusted event-data feeds. But these strengths address infrastructure credibility, not liquidity formation. While institutional adoption will follow infrastructure credibility, liquidity formation will follow market-building capability. Both are required, and the second is where incumbents have the most ground to make up, either by building these capabilities in-house, which means committing to a category of spend that has historically sat outside the FMI operating model, or by partnering with or acquiring platforms that already have them.

This distinction matters because prediction markets will not become a single homogeneous asset class. Some categories will behave like entertainment products. Others will look more like data products. A smaller set could become genuine risk-transfer markets.

Skepticism is understandable: many prediction markets still look speculative, retail-led, and vulnerable to weak governance, manipulation, or poorly defined settlement rules. But that does not make the underlying market structure irrelevant. Properly governed prediction markets solve a specific set of problems that are only partly addressed by existing market infrastructure. They convert dispersed views into continuously updated probabilities, make binary risks tradable, and create market data around events that are otherwise discussed qualitatively.

That is valuable where outcomes are public, time-bound, and easy to verify.

The value proposition is strongest where three conditions hold. First, the event must matter economically. Second, the outcome must be clearly defined and independently observable. Third, there must be enough natural two-way interest to support liquidity.

Where these conditions hold, prediction markets can become useful tools for forecasting, hedging and decision support, with seven core value propositions: better forecasting, observable probabilities, targeted hedging, internal information discovery, distributional expectations, broader retail access, and aggregated market data products.

Exhibit 19: Prediction market value proposition

Problem	How prediction markets help	Value proposition	
Stale or biased forecasts	Convert dispersed views into a continuously updated market signal as new information arrives	Highest where outcomes are public, time-bound and news-sensitive	● ● ●
No direct probability for an event	Turn a yes/no or range outcome into an observable, market-implied likelihood	Strong fit for elections, macro releases, policy decisions and weather bands.	● ● ●
Lack of distributional expectations	Use multiple strikes or ranges to reveal the shape of expectations, not just a point estimate	High for macro, policy and tail-risk monitoring	● ● ●
Hard-to-hedge binary risk	Create contracts whose payoff is linked directly to a defined state of the world	Useful when the contract closely matches the risk exposure	● ● ●
Dispersed private information inside firms	Let employees or experts express conviction directly rather than through filtered escalation	Potentially valuable for launches, demand, delivery dates and sales forecasts	● ● ●
Market-data/sentiment products	Package aggregated odds as an information feed for analysts, media and trading desks	Strong if markets are liquid, deep and well governed	● ● ●
Retail access to defined-risk event exposure	Offer bounded-loss contracts with simple payoffs around clearly specified events	Medium as an access and education proposition; stronger when tied to economic use cases	● ● ●


○○○ Low ●●● High

Prediction markets address seven needs: from stale forecasts and missing event probabilities to hard-to-hedge binary risk and sentiment-data products, with the strongest value where outcomes are public, time-bound, and news-sensitive.

Source: Oliver Wyman analysis

Exhibit 20: Prediction market addressability framework

In number of contracts, billion, rounded



About 10–15% of the roughly 120 billion ETD contracts traded in 2025 is addressable by event contracts; with net-new institutional and retail volume, the prediction-market TAM reaches an estimated 37–63 billion contracts.

Source: Oliver Wyman analysis

The opportunity for prediction markets should be viewed as the next step in listed-derivatives contract innovation. Listed futures and options already translate economic risks into standardized, exchange-traded contracts. Event-risk contracts extend that logic to outcomes that are currently difficult to trade directly, such as macro releases, central-bank decisions, policy outcomes, weather events, corporate actions, and regulated asset-market events.


The market opportunity for prediction markets will be partially supported by a transition of activity from listed derivatives into prediction markets, particularly, where existing futures and options are currently used as imperfect hedges for event-specific risks. These contracts may evolve into more targeted listed structures, such as those around central-bank decisions, inflation prints, index inclusion, weather outcomes, or corporate events.

This means they would compete with or complement existing ETD volumes. We broadly assess that 10-15% of current ETD volume is addressable by prediction markets. The larger part of FMI-relevant prediction market volume will come from net new volume, where market participants can trade or hedge risks for which no clean listed instrument currently exists today. Two groups are driving the net new volume: institutional participant hedging, speculating and market making, and secondly net new retail penetration, including from crypto-native investors.

Adoption will remain selective though. In the base case, the opportunity is driven by a focused set of event-risk contracts rather than broad conversion of the listed-derivatives market. The strategic question for FMIs is to identify, design and scale the right contracts before liquidity, data, and market-implied probability advantages lock in elsewhere.

Exhibit 21: Prediction market revenue potential by adoption scenario

In \$ billion, annual revenue



Prediction markets could generate \$0.5–1 billion of FMI revenue under limited adoption, \$3–5 billion in a base case, and more than \$13–19 billion under widespread adoption, split across trading, post-trade, data, and technology.

Source: Oliver Wyman analysis

Exhibit 21 lays out three adoption scenarios, with different revenue mixes dependent upon whether prediction markets remain retail-led or become institutional markets.

- **Limited adoption:** only a narrow set of contracts forms liquidity, mostly in retail-led or episodic categories. Revenues remain concentrated in trading fees, with limited need for institutional clearing, settlement, data, or workflow infrastructure. Annual revenues reach only half a billion to one billion dollars.
- **Base case:** While adoption remains selective, a focused set of FMI-addressable contracts are scaling across macro, policy, weather, corporate and regulated asset-market events. The revenue pool reaches between \$3 billion and \$5 billion annually, with 60-70% net new to the industry from new institutional speculation and hedging, plus deeper retail and crypto-native participation. The remainder resembles traditional listed-derivatives contract evolution and competes with adjacent ETD revenue pools. As institutional participation grows, the revenue mix broadens beyond trading: roughly 50% of revenues sit outside trading across post-trade, data, and technology services.

- **Widespread adoption:** a broader set of event-risk contracts becomes liquid and institutionally accepted, supported by regulatory clarity, objective settlement, central clearing, surveillance, trusted event data, and workflow connectivity. The revenue pool could reach somewhere between \$13 billion to \$19 billion annually, with roughly 50% generated outside trading across post-trade, data, and technology services. At this point, prediction markets evolve from a trading-fee opportunity into a broader event-risk infrastructure layer.

For FMIs, the opportunity is to make event-risk markets institutional enough to scale rulebooks, governance, neutral infrastructure, trusted data, clearing, surveillance, and distribution. Winners will build the rails around these markets before liquidity and data advantages lock in elsewhere, rather than competing only on front-end experience. That requires closing the market-building gap with crypto and DeFi-native platforms, while compressing product-launch cycles, so new contracts can form liquidity faster.

Actions for FMI management

- Build institutional-grade rails for most economically relevant event-risk contracts
- Own the data layer for new markets
- Build market-building muscle and compress product launch cycles

3C. Compute as an asset class

With the rise of AI, a further growth theme has moved increasingly to the centre: Compute as an asset class.

The market for computing capacity ("Compute") now displays many of the fundamental building blocks of a financial market. The development of the market for compute shows the early hallmarks of commodity financialization.

Compute has the fundamental building blocks of a financial market, with its early days bearing the classic hallmarks of financializing commodities.

First, Compute has become a critical economic input. Global data centre capacity is expected to almost double between 2025 and 2030, from approximately 75 GW to around 140 GW, driven by AI workloads and continued cloud penetration. Analysts place the expected investment in data, technology and infrastructure over this period at over \$5 trillion.

Second, uneven demand growth and supply scarcity are creating persistent volatility and market risk. Supply constraints exist at multiple levels, from bottlenecks in GPU and component availability, to limitations around data centre power access, grid connection, infrastructure permitting and development lead times. These constraints create the conditions for price dispersion, forward-market formation and demand for risk-transfer.

Third, a physical marketplace already exists. End-users, including AI companies, enterprises, model developers and inference providers, procure compute from hyperscalers, neoclouds and specialist infrastructure providers, often on a dollar-per-GPU-hour basis. This market is also evolving beyond simple on-demand usage. Customers increasingly balance spot access, reserved capacity and forward commitments, weighing flexibility against certainty of price and availability. On-demand procurement offers maximum optionality, but leaves users exposed to both price spikes and capacity shortfalls. Reserved and forward structures provide greater predictability, creating a natural bridge to financial hedging markets.

This gives compute a clear commercial hedging logic. Providers that own or control GPU capacity are structurally exposed to utilization and revenue volatility; users of compute are exposed to future cost volatility and availability risk. A mature compute market would therefore bring together capacity providers, commercial users, banks, liquidity providers and financial investors, with risk transferred between parties that have different natural exposures and time horizons.

Despite this potential, standardization remains the central challenge. Customers face significant variation in pricing across chip models, geographies, providers, contract terms, utilization levels and service quality. More importantly, compute is not only non-standard in price terms; it is non-standard in performance terms. The same nominal unit of compute can produce different outcomes depending on hardware configuration, networking, memory, latency, location, software stack and provider operating standards. This means the infrastructure opportunity is not limited to price discovery. It also requires trusted measurement, performance verification and quality assurance.

Benchmark providers are beginning to address this problem, offering a similar standardization role to that played by Price Reporting Agencies (PRAs) and index providers in other financialized markets. Not unlike many of its predecessors, the challenge of developing a robust compute benchmark must include the definition of standards for hardware configuration, contract type, geography, performance characteristics and other quality factors that result in a reference price that is sufficiently representative of the global market to be useful for hedging purposes. Even then, basis risk will remain. A user hedging a specific requirement in a specific region will not be perfectly matched to a broad benchmark, but the objective is to create a benchmark with sufficiently broad appeal to deliver transparency and liquidity in the market for compute.

Given the inherently non-standard nature and regional variability in the market for compute, a healthy over the counter (OTC) market is likely to persist, supported by brokered liquidity, bespoke forward contracts and potentially by electronic request-for-quote (RFQ) platforms. Listed venues will have an important role where contracts can be standardized and benefit from clearing, margining, distribution and liquidity formation.

These developments create opportunities for FMIs and data providers globally in the order of \$0.6-1 billion in annual revenues. Importantly, these opportunities


may develop in local and regional markets as well as at global FMI groups. While incumbent global FMIs have advantages in listed product distribution, clearing, margin optimization and participant connectivity, the market for compute may develop more locally, not unlike global power markets. Power availability, data centre location, grid constraints, permitting, regulation, latency and AI ecosystem density all drive local pricing differences with a resulting opportunity for regional FMIs to build domestic compute markets at the centre of their national AI ecosystems.

While the initial discussion has been focused on “GPU-hours” there are many other routes in which this nascent asset class could develop. For example, token-price benchmarks could emerge as wholesale procurement of AI consumption becomes more systematized and consumption measures become more standardized. Secondary markets for refurbished GPUs, residual-value curves and chip financing could also develop as providers assess whether to deploy new capacity, redeploy older hardware or finance assets against expected future cash flows. In parallel, forward compute contracts introduce counterparty and performance risk, creating potential demand for collateral, clearing, credit intermediation and eventually compute-linked credit products.

Actions for FMI management

- Shape the market structure and build a targeted regional infrastructure

Exhibit 22: Compute markets value chain and landscape



Compute is emerging as a financializing commodity, opening the door for FMIs to shape its market structure, OTC, listed, and post-trade, and build targeted regional infrastructure.
Source: Oliver Wyman analysis

04

Conclusion

The base case set out in this report rests on a simple premise. Sector beta, which includes secular deepening in capital markets, rising risk complexity, continued electronification, and broader retail participation, remains intact and will underpin around 8% annual growth now through 2030. What is shifting is the alpha.

Enterprise AI is pressuring parts of the data and technology portfolio, while creating opportunity in defensible workflows. Tokenization and prediction markets are expanding the perimeter of market infrastructure for incumbents that move early.

Every senior team in the sector sees these trends. Differentiation in the next cycle will come from how clearly each FMI organizes itself around four strategic positioning choices: market-of-record, system-of-work, and source-of-origin. It will also come from how FMIs lean into changes in the operating model required.

The exhibit below outlines four sets of priority actions needed to achieve success over the next cycle: one for each strategic positioning choice, plus an agenda for the operating model.

Exhibit 23: FMI management agenda through 2030

Market-of-record (infrastructure) • Defend post-trade control points on new rails and build ecosystems • Prioritize tokenization where revenue relevance and benefits are highest • Treat regional and regulatory competition as a first-order strategic issue • Build institutional-grade rails for most economically relevant event risk contracts • Shape the 'compute' market structure and build targeted regional infrastructure	System-of-work (technology) • Embed AI into defensible workflow pools • Reprice AI-enabled offerings toward outcomes • Drive usage uplift through copilots, agents and automation • Build token-aware and event-risk workflow modules	Source-of-origin (data) • Segment data revenue pools by defensibility • Protect and reprice fortress data • Re-design, link or exit exposed pools • Own the data layer for new markets
Operating model • Create a top-down cross-organizational vision for AI • Industrialize AI through a factory model • Convert AI productivity into margin uplift • Build or acquire market-building muscle and compress product launch cycle • Make AI risk and governance scalable		

The management agenda spans four choices: defending the market-of-record (infrastructure), the system-of-work (technology), and the source-of-origin (data), and reshaping the operating model, underpinned by scalable AI risk and governance.

Source: Oliver Wyman analysis

Market-of-record (infrastructure)

Winners will defend trusted trade and post-trade economics as activity moves to new rails and jurisdictions. They also will build the institutional rails for new tokenized asset classes and event-risk markets before liquidity and data advantages lock in elsewhere.

01

Defend control points on new rails and build ecosystems. Secure roles in issuance, registry, clearing, settlement, custody, margining, collateral management, and asset servicing as tokenization shifts parts of today's economics to token-native infrastructure. The largest tokenization opportunities sit in post-trade functions, so the strategic priority and need will be to determine which organizations will be first to operate the trusted control points and build liquidity and ecosystems around them.

02

Prioritize tokenization in areas in which revenue relevance and benefits are highest. Focus first will be on crypto post-trade and derivatives, syndicated loans, MMFs and increased velocity for equities and fixed income: crypto needs institutional clearing, custody, settlement, margining, surveillance, and risk controls. Syndicated loans need digital ownership records, entitlement management, transfers, and servicing; MMFs offer the clearest near-term revenue case through collateral, settlement, and liquidity-management integration.

03

Treat regional and regulatory competition as a first-order issue. Engage early with jurisdictions seeking to protect or repatriate clearing, settlement, and custody-linked revenues. Make sovereign, central bank, and regulatory partnerships part of the core infrastructure strategy rather than a peripheral policy exercise.

04

Build institutional-grade rails for most economically relevant event-risk contracts. Focus on economically relevant categories with objective settlement criteria, such as macro, policy, weather, corporate, and regulated asset-market events. Provide the rulebooks, clearing, settlement, surveillance, and trusted event-data infrastructure required for institutional adoption.

05

Shape the Compute market structure and build targeted regional infrastructure. Focus on the standards that make Compute tradable, including contract units, quality tiers, delivery locations, reference pricing, and settlement conventions. Given Compute's heterogeneity, the market is likely to develop through a mix of OTC/RFQ execution, listed hedging, benchmarks, data, and post-trade services, with strong regional overlays.

System-of-work (technology)

Winners will defend and enhance embedded technology positions where they sit deep in client workflows, combining with proprietary data and repricing for value, and build token-aware and event-risk workflow modules to capture the new markets.

01

Embed AI into defensible workflow pools.

Prioritize AI product investment where FMIs already sit in trusted, mission-critical, or regulated workflows, including surveillance, post-trade exception handling, OMS/EMS decision support, risk management, and portfolio workflows.

02

Reprice AI-enabled offerings toward outcomes.

Shift commercial models away from seats and feeds toward measurable outcomes such as risk reduction, time saved, compliance readiness, workflow automation, and operational resilience. Bundle AI capabilities into premium tiers where client value is demonstrable.

03

Drive usage uplift through copilots, agents, and automation.

Put AI directly into daily client workflows and APIs to increase frequency of use, broaden user personas and expand consumption-based revenues.

04

Build token-aware and event-risk workflow modules.

Develop the workflow layer required to integrate new markets into existing client processes: token identifiers, entitlement records, collateral and settlement workflows, event-risk APIs, probability feeds and connectivity into trading and risk systems.

Source-of-origin (data)

Winners will defend data positions where content is proprietary, authoritative, regulatory-mandated or embedded into client distribution, bundling exposed analytics into solutions anchored on these moats and repricing toward outcomes, and build the trusted reference, pricing and event-data layers for tokenized asset classes and event-risk markets.

01

Segment data pools by defensibility. Classify revenue pools as Fortress, Augment, or Contest, based on proprietary content, authority, regulatory and audit relevance, workflow embeddedness, network effects and trust, and link investment, pricing, bundling and exit decisions to this profile.

02

Protect and reprice Fortress data. Defend indices, benchmarks, authoritative reference data, regulated price formation, depository records, trade reporting, surveillance signals, and audit-trail data through outcome-based pricing. These datasets are uniquely defensible because they derive value from authority, regulation and trust, not just analytical insight.

03

Redesign, link, or exit exposed pools. For generic analytics, shallow reporting, aggregated historic market data and non-proprietary datasets, embed AI where it can defend value, bundle exposed functionality into Fortress offerings where possible, and exit where moats cannot be rebuilt. Inaction risks turning less connected data and workflow extensions into a drag on FMI group value.

04

Own the data layer for new markets. Establish source-of-origin positions in token identifiers, entitlement records, reference data, pricing sources, audit trails, market-implied probabilities, and trusted event-risk data feeds before standards, liquidity, and distribution advantages are set by alternative platforms.

Operating model

Winners will treat AI as an operating model lever. They will look to industrialize AI, including governance, so that it converts cost reduction into EBITDA margin uplift and unlock the AI revenue upside at scale. They will also close the market-building capability gap relative to crypto and DeFi-native platforms, building the distribution, incentive, community, and product-launch capabilities that turn new asset classes contracts into liquid markets.

01

Create a top-down cross-organizational vision for AI. Combine top-down and bottom-up AI efforts, cutting across organizational siloes to leverage the opportunities of agentic AI.

02

Industrialize AI through a factory model. Build reusable components, shared data and feature layers, model governance, risk controls, and a value control tower. Sequence deployment from corporate and enablement functions into regulated functions only where explainability, auditability, resilience and model-risk controls are robust (and the business case stacks up).

03

Convert AI productivity into margin uplift. Track cycle-time reduction, first-time-right rates, cost-to-serve, reuse and realized margin impact through hard KPIs.

04

Build or acquire market-building muscle and compress product launch cycles. Develop distribution, incentive design, community engagement, liquidity formation, and data-as-media capabilities. Use MVP discipline, explicit user feedback loops, smaller initial commitments and clear kill criteria for contracts or products that do not form liquidity.

05

Make AI risk and governance scalable. Establish a top-of-house AI governance framework that enables, rather than slows, responsible AI adoption at scale. As use cases move from internal support tools toward agentic and autonomous applications, FMIs will need continuous monitoring, automated testing and validation, clear line-of-defense accountability, and robust controls for model drift, explainability, cyber resilience, third-party dependency, and operational continuity.

Authors

Magnus Burkl CFA

Partner, Global Head of Financial Infrastructure
magnus.burkl@oliverwyman.com

Hiten Patel

Partner, Global Head of Financial Infrastructure,
Technology and Services
hiten.patel@oliverwyman.com

Alexandre Cailleteau

Partner, Global Head of Data and Information Services
alexandre.cailleteau@oliverwyman.com

Michael Wagner

Partner
michael.wagner@oliverwyman.com

Nikolai Dienerowitz

Partner
nikolai.dienerowitz@oliverwyman.com

Binnie Sacks

Partner
benjamin.sacks@oliverwyman.com

Marcus Hoddinott

Principal
marcus.hoddinott@oliverwyman.com

Daniel Martin

Principal
daniel.martin@oliverwyman.com

Nasser Al-Shamma

Engagement Manager
nasser.alshamma@oliverwyman.com

Ukshin Berisha

Engagement Manager
ukshin.berisha@oliverwyman.com

Contributors

Pete Clarke

Partner, Americas Head
of Financial Infrastructure,
Technology and Services

Dylan Walsh

Partner

Lauren Lyons

Partner

Paul Calvey

Partner

Jasper Yip

Partner

Archie Stebbings

Partner

Ben Dietl

Partner

Gianluca Corradi

Partner

Peter Dykstra

Partner

Michael Ho

Partner

Benjamin Holroyd

Partner

Andreas Lagerqvist

Partner

Rebecca Young

Partner

Katie Sessa

Principal

Angella Jiang

Principal

Shubhangi Sharma

Principal

Andrin Boegli

Principal

Sam Maeglin

Principal

Karl Gridl

Principal

Maeve Koetl

Principal

Celia Yao

Principal

Kennedy Urban

Engagement Manager

Elena Camiciola

Engagement Manager

Thomas Wong

Engagement Manager

Michael Magumbe

FITS Research Lead

Oliver Wyman, a Marsh (NYSE: MRSH) business, is a management consulting firm driven by deep industry insight, bold innovation, and a collaborative approach that cuts through complexity to help organizations navigate their most defining transformative moments. Marsh is a global leader in risk, reinsurance and capital, people and investments, and management consulting, advising clients in 130 countries. With annual revenue of \$27 billion and more than 95,000 colleagues, Marsh helps build the confidence to thrive through the power of perspective.

For more information, visit oliverwyman.com, or follow us on [LinkedIn](#) and [X](#).

Americas
+1 212 541 8100

Europe
+44 20 7333 8333

Asia Pacific
+65 6510 9700

India, Middle East & Africa
+971 (0) 4 425 7000

Copyright 2026 Oliver Wyman

All rights reserved. This report may not be reproduced or redistributed, in whole or in part, without the written permission of Oliver Wyman and Oliver Wyman accepts no liability whatsoever for the actions of third parties in this respect.

The information and opinions in this report were prepared by Oliver Wyman. This report is not investment advice and should not be relied on for such advice or as a substitute for consultation with professional accountants, tax, legal or financial advisors. Oliver Wyman has made every effort to use reliable, up-to-date and comprehensive information and analysis, but all information is provided without warranty of any kind, express or implied. Oliver Wyman disclaims any responsibility to update the information or conclusions in this report. Oliver Wyman accepts no liability for any loss arising from any action taken or refrained from as a result of information contained in this report or any reports or sources of information referred to herein, or for any consequential, special or similar damages even if advised of the possibility of such damages. The report is not an offer to buy or sell securities or a solicitation of an offer to buy or sell securities. This report may not be sold without the written consent of Oliver Wyman.